

# Flowers Recognition

## Introduction

We will be using the [Flowers Recognition](#) database.

This dataset contains 4242 images of flowers. The data collection is based on the data flicr, google images, yandex images.

The pictures are divided into five classes: chamomile, tulip, rose, sunflower, dandelion. For each class there are about 800 photos. Photos are not high resolution, about 320x240 pixels. Photos are not reduced to a single size, they have different proportions!

## Objective

we can use this dataset to recognize plants from the photo usning **Convolutional Neural Networks**.

## EDA

```
In [ ]: class_names=['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']

In [ ]:
# Ignore the warnings
import warnings
warnings.filterwarnings('always')
warnings.filterwarnings('ignore')

# data visualisation and manipulation
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
from random import shuffle

#configure
# sets matplotlib to inline and displays graphs below the corresponding cell.
%matplotlib inline
style.use('fivethirtyeight')
sns.set(style='whitegrid',color_codes=True)

#model selection
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,confusion_matrix
from sklearn.preprocessing import LabelEncoder

#preprocess.
from tensorflow.keras.preprocessing.image import ImageDataGenerator

#dl Libraraies
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

```
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical

# specifically for cnn
import tensorflow as tf
from tensorflow.keras.layers import Flatten, Activation
from tensorflow.keras.layers import Conv2D, MaxPooling2D
```

In [ ]:

```
X=[]
Z=[]
IMG_SIZE=150
FLOWER_DAISY_DIR='flowers/daisy'
FLOWER_SUNFLOWER_DIR='flowers/sunflower'
FLOWER_TULIP_DIR='flowers/tulip'
FLOWER_DANDI_DIR='flowers/dandelion'
FLOWER_ROSE_DIR='flowers/rose'
```

In [ ]:

```
def assign_label(img,flower_type):
    return flower_type
```

In [ ]:

```
def make_train_data(flower_type,DIR):
    for img in tqdm(os.listdir(DIR)):
        label=assign_label(img,flower_type)
        path = os.path.join(DIR,img)
        img = cv2.imread(path,cv2.IMREAD_COLOR)
        img = cv2.resize(img, (IMG_SIZE,IMG_SIZE))

        X.append(np.array(img))
        Z.append(str(label))
```

In [ ]:

```
# specifically for manipulating zipped images and getting numpy arrays of pixel values
import cv2
from tqdm import tqdm
import os
```

In [ ]:

```
make_train_data('Daisy',FLOWER_DAISY_DIR)
print(len(X))
```

```
100%|██████████| 764/764 [00:21<00:00, 35.66it/s]
764
```

In [ ]:

```
make_train_data('Sunflower',FLOWER_SUNFLOWER_DIR)
print(len(X))
```

```
100%|██████████| 733/733 [00:13<00:00, 52.77it/s]
1497
```

In [ ]:

```
make_train_data('Tulip',FLOWER_TULIP_DIR)
print(len(X))
```

```
100%|██████████| 984/984 [00:16<00:00, 59.65it/s]
```

2481

```
In [ ]: make_train_data('Dandelion', FLOWER_DANDI_DIR)
print(len(X))
```

```
100%|██████████| 1052/1052 [00:26<00:00, 39.69it/s]
3533
```

```
In [ ]: make_train_data('Rose', FLOWER_ROSE_DIR)
print(len(X))
```

```
100%|██████████| 784/784 [00:04<00:00, 185.63it/s]
4317
```

```
In [ ]: def display_random_image(class_names, images, labels):
    """
        Display a random image from the images array and its correspond label from t
    """

    index = np.random.randint(len(labels))
    plt.figure()
    plt.imshow(images[index])
    plt.xticks([])
    plt.yticks([])
    plt.grid(True)
    plt.title('Image #{} : {}'.format(index, labels[index]))
    plt.show()
```

```
In [ ]: display_random_image(class_names, X, Z)
```



```
In [ ]: le=LabelEncoder()
Y=le.fit_transform(Z)
Y=to_categorical(Y,5)
X=np.array(X)
X=X/255
```

```
In [ ]: x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.25,random_state=42)
```

```
In [ ]: x_train.shape  
Out[ ]: (3237, 150, 150, 3)
```

```
In [ ]: np.random.seed(42)
```

## Modelling using CNN.

```
In [ ]:  
model = Sequential()  
model.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))  
model.add(MaxPooling2D(pool_size=(2,2)))  
  
model.add(Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',activation ='relu'))  
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))  
  
model.add(Conv2D(filters = 96, kernel_size = (3,3),padding = 'Same',activation ='relu'))  
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))  
  
model.add(Conv2D(filters = 96, kernel_size = (3,3),padding = 'Same',activation ='relu'))  
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))  
  
model.add(Flatten())  
model.add(Dense(512))  
model.add(Activation('relu'))  
model.add(Dense(5, activation = "softmax"))
```

```
In [ ]:  
batch_size=128  
epochs=50  
  
from keras.callbacks import ReduceLROnPlateau  
red_lr= ReduceLROnPlateau(monitor='val_acc',patience=3,verbose=1,factor=0.1)
```

```
In [ ]:  
datagen = ImageDataGenerator(  
    featurewise_center=False, # set input mean to 0 over the dataset  
    samplewise_center=False, # set each sample mean to 0  
    featurewise_std_normalization=False, # divide inputs by std of the dataset  
    samplewise_std_normalization=False, # divide each input by its std  
    zca_whitening=False, # apply ZCA whitening  
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180  
    zoom_range = 0.1, # Randomly zoom image  
    width_shift_range=0.2, # randomly shift images horizontally (fraction of total width)  
    height_shift_range=0.2, # randomly shift images vertically (fraction of total height)  
    horizontal_flip=True, # randomly flip images  
    vertical_flip=False) # randomly flip images  
  
datagen.fit(x_train)
```

```
In [ ]: model.compile(optimizer=Adam(lr=0.001),loss='categorical_crossentropy',metrics=['acc'])
```

```
In [ ]: model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 150, 150, 32)	2432
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 64)	0
conv2d_2 (Conv2D)	(None, 37, 37, 96)	55392
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 96)	0
conv2d_3 (Conv2D)	(None, 18, 18, 96)	83040
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 96)	0
flatten (Flatten)	(None, 7776)	0
dense (Dense)	(None, 512)	3981824
activation (Activation)	(None, 512)	0
dense_1 (Dense)	(None, 5)	2565
<hr/>		
Total params: 4,143,749		
Trainable params: 4,143,749		
Non-trainable params: 0		

In [ ]:

```
History = model.fit_generator(datagen.flow(x_train,y_train, batch_size=batch_size),
                               epochs = epochs, validation_data = (x_test,y_test),
                               verbose = 1, steps_per_epoch=x_train.shape[0] // batch
```

```
Epoch 1/50
25/25 [=====] - 62s 1s/step - loss: 1.4996 - accuracy: 0.33
77 - val_loss: 1.1777 - val_accuracy: 0.5176
Epoch 2/50
25/25 [=====] - 20s 784ms/step - loss: 1.2336 - accuracy: 0.4693
0 - val_loss: 1.0705 - val_accuracy: 0.5880
Epoch 3/50
25/25 [=====] - 19s 761ms/step - loss: 1.1225 - accuracy: 0.5429
0 - val_loss: 1.0079 - val_accuracy: 0.5824
Epoch 4/50
25/25 [=====] - 19s 738ms/step - loss: 1.0591 - accuracy: 0.5764
0 - val_loss: 1.0275 - val_accuracy: 0.5769
Epoch 5/50
25/25 [=====] - 22s 852ms/step - loss: 0.9921 - accuracy: 0.6092
0 - val_loss: 0.9213 - val_accuracy: 0.6389
Epoch 6/50
25/25 [=====] - 19s 777ms/step - loss: 0.9332 - accuracy: 0.6375
0 - val_loss: 0.8440 - val_accuracy: 0.6639
Epoch 7/50
25/25 [=====] - 20s 784ms/step - loss: 0.8963 - accuracy: 0.6497
0 - val_loss: 0.8979 - val_accuracy: 0.6509
Epoch 8/50
25/25 [=====] - 20s 777ms/step - loss: 0.9035 - accuracy: 0.6571
0 - val_loss: 0.8089 - val_accuracy: 0.6861
Epoch 9/50
25/25 [=====] - 19s 760ms/step - loss: 0.8623 - accuracy:
```

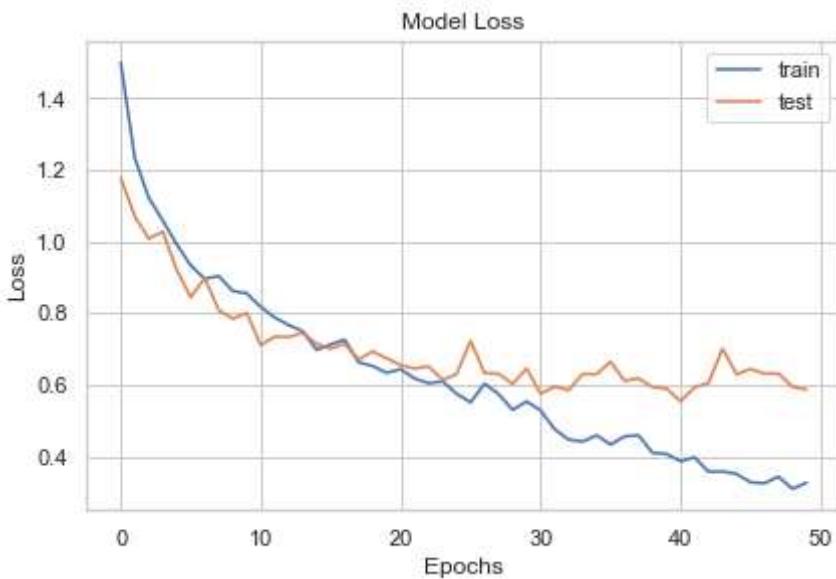
0.6748 - val\_loss: 0.7841 - val\_accuracy: 0.6954  
Epoch 10/50  
25/25 [=====] - 21s 817ms/step - loss: 0.8556 - accuracy:  
0.6652 - val\_loss: 0.8014 - val\_accuracy: 0.6722  
Epoch 11/50  
25/25 [=====] - 20s 804ms/step - loss: 0.8181 - accuracy:  
0.6850 - val\_loss: 0.7105 - val\_accuracy: 0.7222  
Epoch 12/50  
25/25 [=====] - 19s 732ms/step - loss: 0.7881 - accuracy:  
0.7002 - val\_loss: 0.7345 - val\_accuracy: 0.7120  
Epoch 13/50  
25/25 [=====] - 18s 723ms/step - loss: 0.7669 - accuracy:  
0.7134 - val\_loss: 0.7337 - val\_accuracy: 0.7185  
Epoch 14/50  
25/25 [=====] - 20s 789ms/step - loss: 0.7495 - accuracy:  
0.7186 - val\_loss: 0.7451 - val\_accuracy: 0.7157  
Epoch 15/50  
25/25 [=====] - 19s 752ms/step - loss: 0.6975 - accuracy:  
0.7414 - val\_loss: 0.7143 - val\_accuracy: 0.7204  
Epoch 16/50  
25/25 [=====] - 19s 740ms/step - loss: 0.7130 - accuracy:  
0.7343 - val\_loss: 0.7011 - val\_accuracy: 0.7167  
Epoch 17/50  
25/25 [=====] - 19s 733ms/step - loss: 0.7262 - accuracy:  
0.7289 - val\_loss: 0.7147 - val\_accuracy: 0.7278  
Epoch 18/50  
25/25 [=====] - 21s 831ms/step - loss: 0.6629 - accuracy:  
0.7444 - val\_loss: 0.6708 - val\_accuracy: 0.7407  
Epoch 19/50  
25/25 [=====] - 21s 808ms/step - loss: 0.6527 - accuracy:  
0.7491 - val\_loss: 0.6936 - val\_accuracy: 0.7370  
Epoch 20/50  
25/25 [=====] - 21s 817ms/step - loss: 0.6338 - accuracy:  
0.7655 - val\_loss: 0.6740 - val\_accuracy: 0.7389  
Epoch 21/50  
25/25 [=====] - 35s 1s/step - loss: 0.6439 - accuracy: 0.76  
23 - val\_loss: 0.6551 - val\_accuracy: 0.7481  
Epoch 22/50  
25/25 [=====] - 21s 811ms/step - loss: 0.6180 - accuracy:  
0.7642 - val\_loss: 0.6456 - val\_accuracy: 0.7620  
Epoch 23/50  
25/25 [=====] - 18s 723ms/step - loss: 0.6047 - accuracy:  
0.7694 - val\_loss: 0.6521 - val\_accuracy: 0.7454  
Epoch 24/50  
25/25 [=====] - 19s 735ms/step - loss: 0.6096 - accuracy:  
0.7629 - val\_loss: 0.6138 - val\_accuracy: 0.7667  
Epoch 25/50  
25/25 [=====] - 21s 809ms/step - loss: 0.5752 - accuracy:  
0.7765 - val\_loss: 0.6300 - val\_accuracy: 0.7528  
Epoch 26/50  
25/25 [=====] - 21s 842ms/step - loss: 0.5512 - accuracy:  
0.7974 - val\_loss: 0.7228 - val\_accuracy: 0.7231  
Epoch 27/50  
25/25 [=====] - 21s 847ms/step - loss: 0.6035 - accuracy:  
0.7626 - val\_loss: 0.6327 - val\_accuracy: 0.7583  
Epoch 28/50  
25/25 [=====] - 21s 817ms/step - loss: 0.5740 - accuracy:  
0.7806 - val\_loss: 0.6308 - val\_accuracy: 0.7491  
Epoch 29/50  
25/25 [=====] - 21s 819ms/step - loss: 0.5303 - accuracy:  
0.7999 - val\_loss: 0.6022 - val\_accuracy: 0.7630  
Epoch 30/50  
25/25 [=====] - 21s 830ms/step - loss: 0.5537 - accuracy:  
0.7858 - val\_loss: 0.6458 - val\_accuracy: 0.7546

```
Epoch 31/50
25/25 [=====] - 21s 840ms/step - loss: 0.5285 - accuracy: 0.8077 - val_loss: 0.5754 - val_accuracy: 0.7944
Epoch 32/50
25/25 [=====] - 20s 790ms/step - loss: 0.4765 - accuracy: 0.8327 - val_loss: 0.5954 - val_accuracy: 0.7889
Epoch 33/50
25/25 [=====] - 21s 835ms/step - loss: 0.4475 - accuracy: 0.8340 - val_loss: 0.5851 - val_accuracy: 0.7861
Epoch 34/50
25/25 [=====] - 21s 815ms/step - loss: 0.4416 - accuracy: 0.8382 - val_loss: 0.6304 - val_accuracy: 0.7722
Epoch 35/50
25/25 [=====] - 21s 820ms/step - loss: 0.4592 - accuracy: 0.8241 - val_loss: 0.6300 - val_accuracy: 0.7787
Epoch 36/50
25/25 [=====] - 21s 839ms/step - loss: 0.4333 - accuracy: 0.8434 - val_loss: 0.6643 - val_accuracy: 0.7630
Epoch 37/50
25/25 [=====] - 21s 825ms/step - loss: 0.4562 - accuracy: 0.8260 - val_loss: 0.6115 - val_accuracy: 0.7806
Epoch 38/50
25/25 [=====] - 20s 778ms/step - loss: 0.4592 - accuracy: 0.8308 - val_loss: 0.6183 - val_accuracy: 0.7731
Epoch 39/50
25/25 [=====] - 20s 798ms/step - loss: 0.4106 - accuracy: 0.8479 - val_loss: 0.5936 - val_accuracy: 0.7870
Epoch 40/50
25/25 [=====] - 20s 772ms/step - loss: 0.4074 - accuracy: 0.8479 - val_loss: 0.5896 - val_accuracy: 0.7926
Epoch 41/50
25/25 [=====] - 20s 790ms/step - loss: 0.3868 - accuracy: 0.8530 - val_loss: 0.5540 - val_accuracy: 0.7981
Epoch 42/50
25/25 [=====] - 19s 735ms/step - loss: 0.3977 - accuracy: 0.8488 - val_loss: 0.5932 - val_accuracy: 0.7898
Epoch 43/50
25/25 [=====] - 19s 731ms/step - loss: 0.3577 - accuracy: 0.8614 - val_loss: 0.6054 - val_accuracy: 0.7833
Epoch 44/50
25/25 [=====] - 19s 754ms/step - loss: 0.3586 - accuracy: 0.8678 - val_loss: 0.7009 - val_accuracy: 0.7648
Epoch 45/50
25/25 [=====] - 20s 789ms/step - loss: 0.3516 - accuracy: 0.8688 - val_loss: 0.6295 - val_accuracy: 0.7639
Epoch 46/50
25/25 [=====] - 19s 733ms/step - loss: 0.3280 - accuracy: 0.8758 - val_loss: 0.6442 - val_accuracy: 0.7806
Epoch 47/50
25/25 [=====] - 20s 809ms/step - loss: 0.3258 - accuracy: 0.8797 - val_loss: 0.6315 - val_accuracy: 0.7815
Epoch 48/50
25/25 [=====] - 21s 809ms/step - loss: 0.3437 - accuracy: 0.8668 - val_loss: 0.6316 - val_accuracy: 0.7880
Epoch 49/50
25/25 [=====] - 19s 743ms/step - loss: 0.3098 - accuracy: 0.8810 - val_loss: 0.5948 - val_accuracy: 0.8000
Epoch 50/50
25/25 [=====] - 21s 813ms/step - loss: 0.3268 - accuracy: 0.8716 - val_loss: 0.5869 - val_accuracy: 0.8046
```

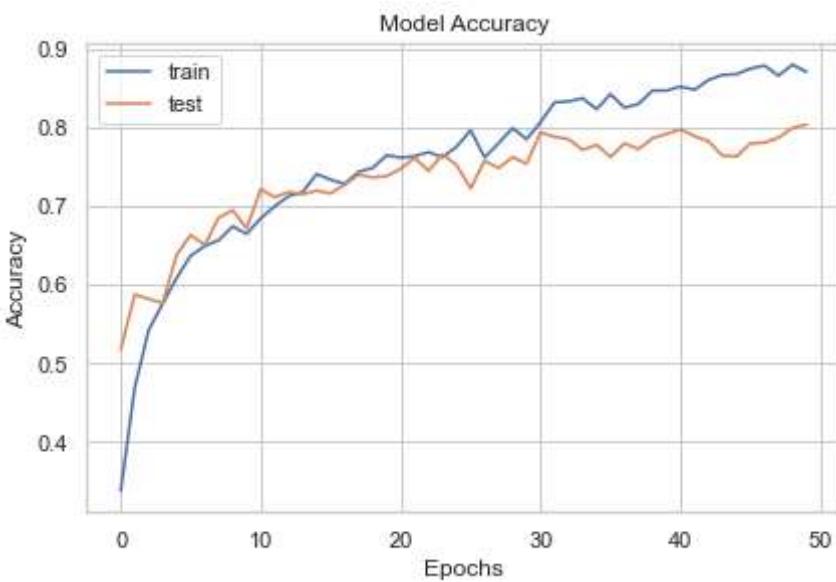
In [ ]:

```
plt.plot(History.history['loss'])
plt.plot(History.history['val_loss'])
```

```
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['train', 'test'])
plt.show()
```



```
In [ ]: plt.plot(History.history['accuracy'])
plt.plot(History.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['train', 'test'])
plt.show()
```



```
In [ ]: # getting predictions on val set.
pred=model.predict(x_test)
pred_digits=np.argmax(pred,axis=1)
```

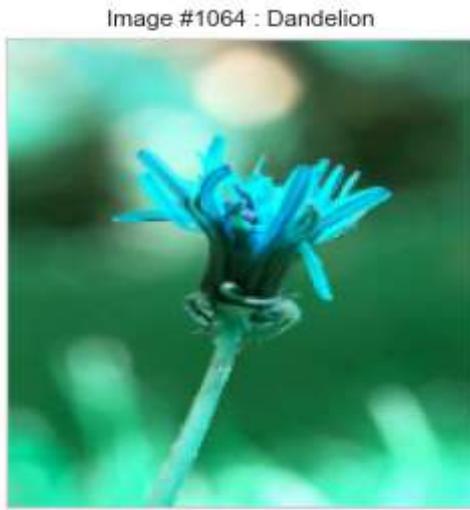
```
In [ ]: test_digits = np.argmax(y_test, axis=1)
```

```
In [ ]: print("Accuracy : {}".format(accuracy_score(test_digits, pred_digits)))
```

```
Accuracy : 0.8046296296296296
```

```
In [ ]: pred_labels = le.inverse_transform(pred_digits)
```

```
In [ ]: display_random_image(class_names, x_test, pred_labels)
```



## Error analysis

```
In [ ]: def display_examples(class_names, images, labels):
    """
        Display 25 images from the images array with its corresponding labels
    """

    fig = plt.figure(figsize=(10,10))
    fig.suptitle("Some examples of images of the dataset", fontsize=16)
    for i in range(25):
        plt.subplot(5,5,i+1)
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)
        plt.imshow(images[i], cmap=plt.cm.binary)
        plt.xlabel(class_names[labels[i]])
    plt.show()
```

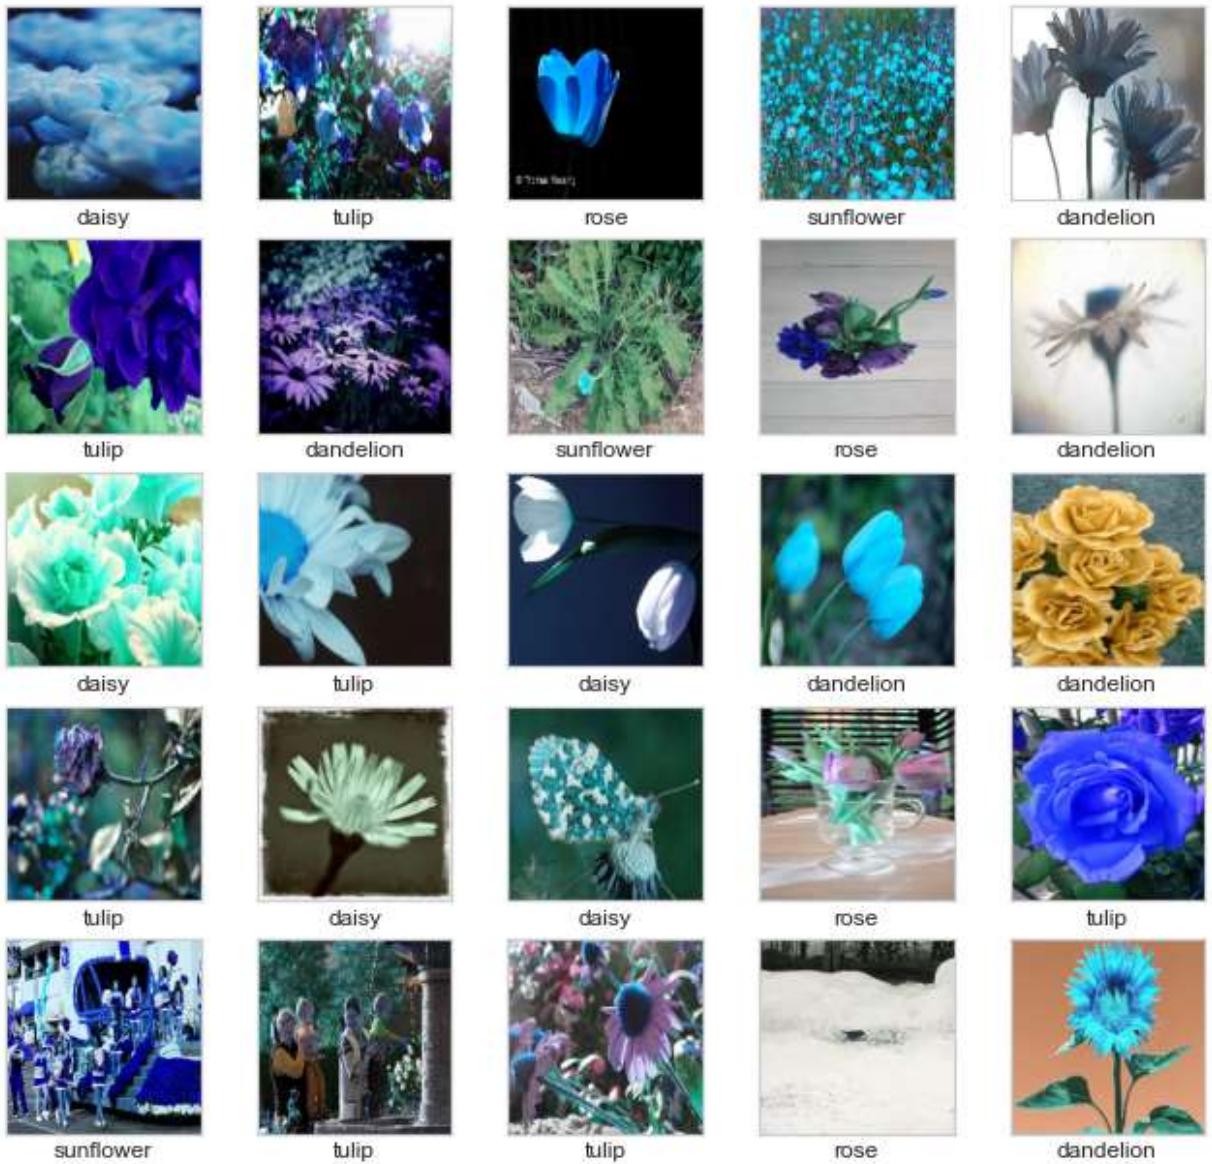
```
In [ ]: def print_mislabeled_images(class_names, test_images, test_labels, pred_labels):
    """
        Print 25 examples of mislabeled images by the classifier, e.g when test_labe
    """

    B00 = (test_labels == pred_labels)
    mislabeled_indices = np.where(B00 == 0)
    mislabeled_images = test_images[mislabeled_indices]
    mislabeled_labels = pred_labels[mislabeled_indices]

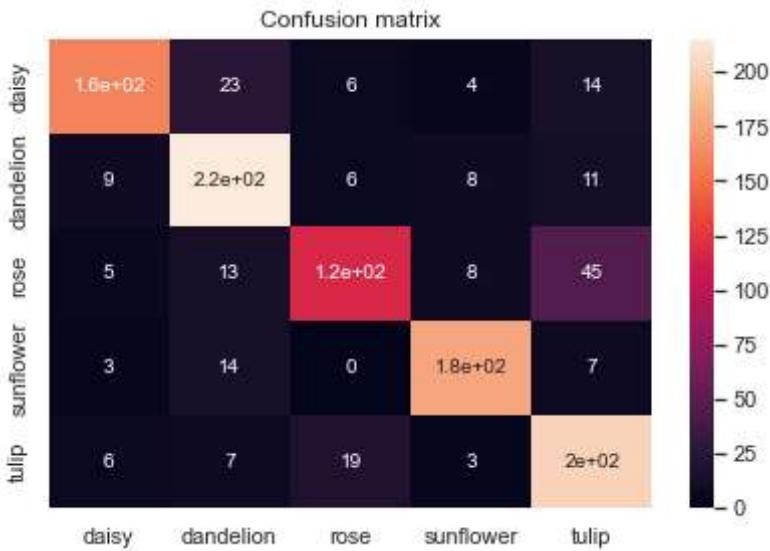
    title = "Some examples of mislabeled images by the classifier:"
    display_examples(class_names, mislabeled_images, mislabeled_labels)
```

```
In [ ]: print_mislabeled_images(class_names, x_test, test_digits, pred_digits)
```

## Some examples of images of the dataset



```
In [ ]:  
CM = confusion_matrix(test_digits, pred_digits)  
ax = plt.axes()  
sns.heatmap(CM, annot=True,  
            annot_kws={"size": 10},  
            xticklabels=class_names,  
            yticklabels=class_names, ax = ax)  
ax.set_title('Confusion matrix')  
plt.show()
```



## Conclusion: The classifier has trouble with 1 kind of images.

It has trouble with `rose` and `tulip`. It can detect `dandelion` very accurately!

## Possible flaws and next steps

We have overfitting problem and we can face it by :

1. using more data
2. regularization(L2, dropout, Data\_augmentation)
3. find better neural network Architecture

or we can stop at 20 epochs as shown in the model complexity curve